



Bayesian analysis for Data Science

Externado para la vida

Mathematics Department



Outline

What is PyMC?

Installing PyMC

Parallel Execution

Basic Example

Hierarchical Example

Diagnostics with ArviZ

Tips and Best Practices

Checklist

Conclusion



What is PyMC?

- PyMC is a probabilistic programming library for Bayesian statistical modeling and inference.
- It uses Theano / Aesara backend for automatic differentiation.
- Enables MCMC, variational inference, and model diagnostics.
- Models are defined using Python syntax.

Installing PyMC

Recommended with conda (Windows/Mac/Linux):

- `conda create -n pymc-env python=3.10`
- `conda activate pymc-env`
- `conda install -c conda-forge pymc`

With pip (less stable):

- `pip install pymc`

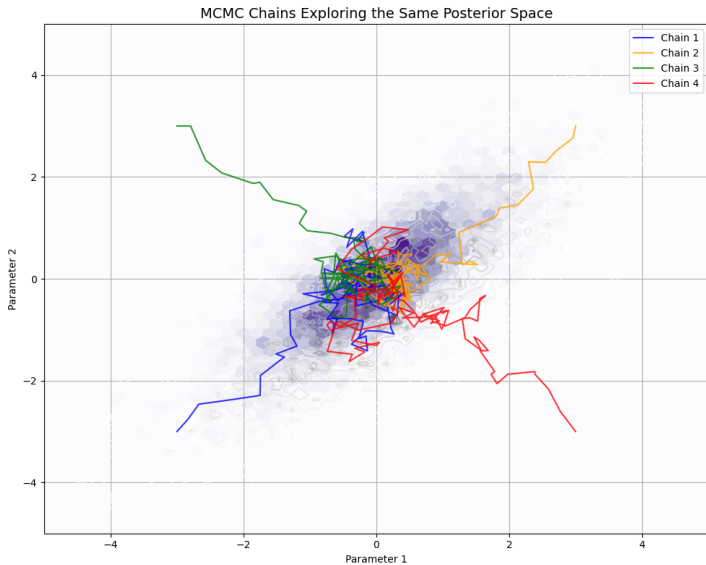
Optional: install Jupyter

- `pip install notebook jupyterlab`

- PyMC uses multiprocessing by default for parallel chains.
- Number of cores can be controlled via `pm.sample()`'s `cores` argument:
- `pm.sample(draws=1000, chains=4, cores=4)`
- Windows needs special care due to multiprocessing spawning.
- Wrap your model inside a function and guard execution:

```
with model(): trace = pm.sample(draws=1000, chains=4,  
cores=4)
```

Chains



Beta-Binomial Example

```
with pm.Model() as beta_binom:
    alpha = pm.HalfNormal("alpha", 10)
    beta = pm.HalfNormal("beta", 10)
    theta = pm.Beta("theta", alpha=alpha, beta=beta)
    y_obs = pm.Binomial("y_obs", n=10, p=theta, observed=7)
    trace = pm.sample(1000)
    pm.model_to_graphviz(beta_binom)
```

Bayesian Linear Regression

```
with pm.Model() as model:  
    alpha = pm.Normal("alpha", mu=0, sigma=10)  
    beta = pm.Normal("beta", mu=0, sigma=10)  
    sigma = pm.HalfNormal("sigma", sigma=1)  
    mu = alpha + beta * x  
    y_obs = pm.Normal("y_obs", mu=mu, sigma=sigma, observed=y)  
    trace = pm.sample()
```


Hierarchical Intercept Model

```
with pm.Model() as model:
    mu_a = pm.Normal("mu_alpha", 0, 5)
    sigma_a = pm.HalfNormal("sigma_alpha", 1)
    alpha = pm.Normal("alpha", mu=mu_a, sigma=sigma_a, shape=n)
    beta = pm.Normal("beta", mu=0, sigma=1)
    mu = alpha[group_idx] + beta * x
    y_obs = pm.Normal("y_obs", mu=mu, sigma=1, observed=y)
    trace = pm.sample()
```

Use ArviZ to check convergence and summarize:

- `az.plot_trace(trace)`
- `az.summary(trace)`
- `az.plot_forest(trace)`
- `az.plot_posterior(trace)`

Tips for Modeling

- Use weakly informative priors unless you have strong prior knowledge.
- Standardize predictors if needed.
- Start with simple models, then add complexity.
- Check \hat{R} and effective sample size.
- Use `idata = pm.sample(return_inferencedata = True)`

Checklist for Bayesian Models with MCMC

Data Preparation

- Clean and explore the data.
- Simulate data with known parameters.

Model Definition

- Define structure and choose priors.
- Check prior assumptions visually.

Model Construction

- Implement likelihood.

Sampling

- Choose MCMC algorithm, set parameters, monitor warnings.

Checklist Continued

Evaluation

- Trace plots, R-hat, ESS, divergence analysis.

Model Comparison

- WAIC, LOO, Bayes Factor, posterior predictive checks.

Tuning and Validation

- Adjust priors, validate on test set.

Documentation

- Assumptions, interpretation, reproducibility.

Conclusion

- PyMC enables flexible, interpretable Bayesian inference.
- It integrates well with modern data science workflows.
- Combine it with ArviZ for diagnostics and visualization.
- Great for hierarchical models, decision making, and uncertainty quantification.